

SQL Analyzing Business Metrics

At the heart of every great business decision is data. Since most businesses store critical data in SQL databases, a deep understanding of SQL is a necessary skill for every data analyst.

Chief among data analysis tasks is *data aggregation*, the grouping of data to express in summary form. We'll be working with **SpeedySpoon**, a meal delivery app. The folks at SpeedySpoon have asked us to look into their deliveries and help them optimize their process.

This course was developed in partnership with our good friends at [Periscope Data](#).

```
select *  
from orders  
order by id  
limit 100; // select only 100 rows from orders
```

Note that the order and limit clauses **keep the data organized**.

To get the Daily Metrics we need the *date*. Most dates in databases are *timestamps*, which have hours and minutes, as well as the year, month, and day. Timestamps look like 2015-01-05 14:43:31, while dates are the first part: 2015-01-05.

We can easily select the date an item was ordered at with the date function and the ordered_at field:

```
select date(ordered_at)  
from orders; //ordered_at is a col in the orders table and we apply to it the  
date() function
```

Let's get a Daily Count of orders from the orders table. Complete the query using the date function to cast the timestamps in ordered_at to dates.

```
select /**/  
from orders  
order by 1  
limit 100;
```

The **order by 1** statement is a shortcut for **order by date(ordered_at)**. The **1** refers to the first column.

```
select date(ordered_at)  
from orders  
order by 1
```

```
limit 100;
```

enerally, when we count all records in a table we run a query with the count function, as follows:

```
select count(1)
from users;
```

This will treat all rows as a single group, and return one row in the result set - the total count.

To count orders by their dates, we'll use the date and count functions and pair them with the group by clause. Together this will count the records in the table, grouped by date.

For example, to see the user records counted by the date created, we use the date and count functions and group by clause as follows:

```
select date(created_at), count(1)
from users
group by date(created_at)
```

Use the date and count functions and group by clause to count and group the orders by the dates they were ordered_at.

```
select /**/
from orders
group by 1
order by 1;
```

We now have the daily order count!

The order by 1 and group by 1 statements are shortcuts for order by date(ordered_at) and group by date(ordered_at).

```
select date(ordered_at), count(1)
from orders
group by 1
order by 1;
```

We can make a few changes to our Daily Count query to get the revenue.

First, instead of using count(1) to count the rows per date, we'll use round(sum(amount_paid), 2) to add up the revenue per date. Complete the query by adding revenue per date.

Second, we need to join in the order_items table because that table has an amount_paid column representing revenue. Complete the query by adding a join clause where orders.id = order_items.order_id.

```
select date(ordered_at), round(sum(amount_paid),2)
from orders
join order_items on orders.id = order_items.order_id
group by 1
order by 1;
```

Note that the round function rounds decimals to digits, based on the number passed in. Here round(..., 2) rounds the sum paid to two digits.

Nice. Now with a small change, we can find out how much we're making per day for any single dish. What's the daily revenue from customers ordering kale smoothies?

Complete the query by using a where clause to filter the daily sums down to orders where the name = 'kale-smoothie'.

```
select date(ordered_at), round(sum(amount_paid),2)
from orders
join order_items on orders.id = order_items.order_id where order_items.name
= 'kale-smoothie'
group by 1
order by 1;
```

To get the percent of revenue that each item represents, we need to know the total revenue of each item. We will later divide the per-item total with the overall total revenue.

The following query groups and sum the products by price to get the total revenue for each item

```
select name,round(sum(amount_paid),2)
from order_items
group by name
order by 2 desc;
```

Grouping with Case Statements

To see if our smoothie suspicion has merit, let's look at purchases by category. We can group the order items by what type of food they are, and go from there. Since our order_items table does not include categories already, we'll need to make some!

Previously we've been using group by with a column (like order_items.name) or a function (like date(orders.ordered_at)).

We can also use group by with expressions. In this case a case statement is just what we need to build our own categories. case statements are similar to if/else in other languages.

Here's the basic structure of a case statement:

```
case {condition}
  when {value1} then {result1}
  when {value2} then {result2}
  else {result3}
end
```

We'll build our own categories using a case statement. Complete the query below with a case condition of name that lists out each product, and decides its group.

```
select *,
case (order_items.name)
  when 'kale-smoothie' then 'smoothie'
  when 'banana-smoothie' then 'smoothie'
  when 'orange-juice' then 'drink'
  when 'soda' then 'drink'
  when 'blt' then 'sandwich'
  when 'grilled-cheese' then 'sandwich'
  when 'tikka-masala' then 'dinner'
  when 'chicken-parm' then 'dinner'
  else 'other'
end as category
from order_items
order by id
limit 100;
```

id	order_id	name	amount_paid	category
0	114	chicken-parm	10.25	dinner
1	2825	tikka-masala	11.25	dinner
2	4717	blt	9.25	sandwich
3	1035	orange-juice	5.25	drink
4	1023	grilled-cheese	9.25	sandwich
5	4359	chicken-parm	10.25	dinner

2.

Complete the query by using the category column created by the case statement in our previous revenue percent calculation. Add the denominator that will sum the amount_paid.

```
select
case name
  when 'kale-smoothie' then 'smoothie'
  when 'banana-smoothie' then 'smoothie'
  when 'orange-juice' then 'drink'
```

```

when 'soda'          then 'drink'
when 'blt'           then 'sandwich'
when 'grilled-cheese' then 'sandwich'
when 'tikka-masala'   then 'dinner'
when 'chicken-parm'   then 'dinner'
else 'other'
end as category, round(1.0 * sum(amount_paid) /
  (select sum(amount_paid) from order_items) * 100, 2) as pct
from order_items
group by 1 //grouped by category and ordered by pct
order by 2 desc;

```

Here 1.0 * is a shortcut to ensure the database represents the percent as a decimal.

Reorder Rates

While we do know that kale smoothies (and drinks overall) are not driving a lot of revenue, we don't know why. A big part of data analysis is implementing your own metrics to get information out of the piles of data in your database.

In our case, the reason could be that no one likes kale, but it could be something else entirely. To find out, we'll create a metric called *reorder rate* and see how that compares to the other products at SpeedySpoon.

We'll define *reorder rate* as the ratio of the total number of orders to the number of people making those orders. A lower ratio means most of the orders are reorders. A higher ratio means more of the orders are first purchases.

Let's calculate the reorder ratio for all of SpeedySpoon's products and take a look at the results. Counting the total orders per product is straightforward. We count the distinct order_ids in the order_items table.

Complete the query by passing in the distinct keyword and the order_id column name into the count function

```

select name, count(distinct order_id)
from order_items
group by 1
order by 1;

```

Here's a [hint](#) on how to use the count function to count distinct columns in a table.

Now we need the number of people making these orders.

To get that information, we need to join in the orders table and count unique values in the delivered_to field, and sort by the reorder_rate.

Complete the query below. The numerator should count the distinct order_ids. The denominator should count the distinct values of the orders table's delivered_to field (orders.delivered_to).

```
select name, round(1.0 * count(distinct order_items.order_id) /  
    count(distinct orders.delivered_to), 2) as reorder_rate  
from order_items  
join orders on
```

Query Results	
name	reorder_rate
banana-smoothie	10.5
kale-smoothie	7.2
tikka-masala	3.32
grilled-cheese	3.06
chicken-parm	2.72
blt	2.49
soda	2.38
orange-juice	1.5
cake	1.4
Database Schema	
orders4999 rows	
id	int
ordered_at	text
delivered_at	text
delivered_to	int
order_items20000 rows	
id	int
order_id	int
name	text
amount_paid	real

orders.id = order_items.order_id
group by 1

order by 2 desc;

Conclusion

Wow! That's unexpected. While smoothies aren't making a lot of money for SpeedySpoon, they have a very high reorder rate. That means these smoothie customers are strong repeat customers.

Instead of recommending smoothies be taken off the menu, we should talk to the smoothie customers and see what they like so much about these smoothies. There could be an opportunity here to expand the product line, or get other customers as excited as these kale fanatics. Nice work!

Let's generalize what we've learned so far:

- *Data aggregation* is the grouping of data in summary form.
- *Daily Count* is the count of orders in a day.
- *Daily Revenue Count* is the revenue on orders per day.
- *Product Sum* is the total revenue of a product.
- *Subqueries* can be used to perform complicated calculations and create filtered or aggregate tables on the fly.
- ***Reorder Rate* is the ratio of the total number of orders to the number of people making orders.**